

Layer-wise Analysis of Representation Fine-Tuning under resource constraints: Which layers matter most?

Muhammad Saad Azam
Department of Computer Science
University of Engineering and Technology
Lahore, Pakistan
saad.azam1090@gmail.com

Abstract—Representation Fine-Tuning (ReFT) is a PEFT (Parameter Efficient Fine-Tuning) technique which directly intervenes on the output of the hidden representations, which makes it lighter than other PEFTs like LoRA and full fine-tuning. However, its productivity and efficiency has remained largely underexplored when it comes to sub-1B models within resource constraints. This paper presents a layer-wise analysis of ReFT on three SLMs (Small Language Models), evaluated on four benchmark tasks varying in difficulty. By applying ReFT independently to each transformer layer while all the other hyperparameters remain the same, the contribution of each layer is measured. The analysis reveals that task complexity is most sensitive to layer selection, where simpler tasks like sentiment analysis are more or less layer agnostic while reasoning ones require a rather careful selection of the layer for fine-tuning. Further analysis is done to see the trends between training dynamics, token generation bias, weight norms and their correlation with accuracy scores.

Index Terms—representation fine-tuning, parameter-efficient fine-tuning, small language models, layer-wise analysis, resource-constrained settings, transformer interpretability, ReFT

I. INTRODUCTION

Since the introduction of transformers [1], the landscape has changed to building large-scale language models. Today, Large Language Models (LLMs) are comprised of billions of parameters [2]. With this enormous size comes the challenging task of fine-tuning these models for specific tasks. While full fine-tuning still yields more accurate results, the cost of this is enormous. To address this challenge, the field of parameter-efficient finetuning (PEFT) has gathered considerable attention. These methods enable researchers to experiment on state-of-the-art models with limited computational resources.

Different techniques have been introduced to fine-tune LLMs. Adapter layers insert a small module in a frozen transformer and train it [4]. Prefix tuning modifies the inputs to attentions instead of modifying weights or activations [5]. Hu et al. [3] introduced a Low-Rank Adaptation (LoRA) of Large Language Models, a way to fine-tune a model without fine-tuning all of its parameters. All these methods either change the inputs or modify the weights without requiring full finetuning.

One more highly efficient technique of PEFT is Representation Finetuning (ReFT) for Language Models. In ReFT, the

representations are directly modified instead of weights of the models [6]. Hence, ReFT acts as a drop-in intervention in representations and steer the output of the model towards a specific direction. Furthermore, ReFT trains over less parameters compared to a standard LoRA setup and achieves fairly competent performance over various tasks.

Unlike weight-based adapters which modify model’s knowledge base, ReFT operates on the hidden states. ReFT helps the model learn geometric transformations which in result nudge the direction of the model’s output without altering its memory. This distinction is crucial because it signals the success of these representations depends more on the semantic richness of the representations at a given depth, which can vary depending on the model architecture and the task of finetuning.

However, ReFT finetuning has a large hyperparameter search space and it is sensitive to layer selection [6]. Also, if ReFT is applied to multiple layers, the intervention matrices must be kept active during inference, introducing some inference delay. Furthermore, ReFT has not been thoroughly studied for Small Language Models (SLMs) under resource constraints on a limited dataset. The original paper also showcases how ReFT is susceptible to nature of tasks as well, further expanding the hyperparameter search space. This requires investigation on how the search space can be narrowed down based on the model architecture and the dataset constraints in place.

With this backdrop, it can be highlighted that the efficacy of ReFT is regulated by layer-wise specialization and architectural constraints. Existing work has primarily focused on 7B+ models. The behavior of ReFT in the ‘Low-Resource’ areas is underexplored. The research will focus on the following questions:

- What does the trend of loss looks like for the layers? Do some layers tend to overfit in ReFT while others do not?
- How does the accuracy vary across the layers depending on the specific tasks?
- How productive is ReFT for small language models?
- What structural and performance similarities exist in layer-wise behavior across differing model scales?

The rest of the paper is organized as follows: We first present the existing work done when it comes to ReFT and layer-wise analysis. The next section describes the methodology of the paper, highlighting the models, datasets and the configuration variables. Later, we analyze the findings of the experiments through the lens of gradient norms, loss landscapes, accuracies and token generation biases. And lastly, we discuss the findings and identify the limitations of this paper.

II. LITERATURE REVIEW

A. Representation Finetuning

Representation Finetuning (ReFT) is a parameter-efficient finetuning technique for LLMs [6]. It trains interventions over the hidden states to steer model behavior into the direction of the finetuning task at inference time. Unlike other methods that target weights, ReFT modifies the hidden representations. The main inspiration of this methodology is the study of causal interventions to direct the behavior of the model [7]. There are two types of interventions introduced originally: LoReFT and DiReFT. LoReFT learns intervention in linear subspace spanned by a low-rank projection matrix, while DiReFT simply shifts the representation by a fixed vector. DiReFT is more straightforward and efficient but is not as accurate as LoReFT is. LoReFT transformation looks like this:

$$\Phi(h) = h + R^T(Wh + b - Rh)$$

Where h is the original hidden representation, R is the low-rank projection matrix, W is the learned weight matrix and b is the learned bias vector.

B. Linearly Decodable Features

Research has shown that many features encoded in the hidden representations are linearly decodable, simulated with linear probing experiments [8]. This is the area of interpretation that ReFT targets. When features are linearly decodable, a learned linear projection over hidden representations can be quite effective. However, this is not always the case, as other research argues the opposite [9]. This makes ReFT potentially less effective where the relevant subspace cannot be manipulated through linear interventions, varying its efficiency from task to task.

C. Model and Layer-wise Analysis

Much work has shown how layers can be specialized for certain tasks depending on the depth. Song et al., have shown how even the later layers are important for certain generation-based tasks and pruning them can degrade the performance of the LLMs within certain context and evaluation metric [10]. The original ReFT paper also identifies the large parameter space and how targetting different layers can vary the efficiency of the LLM over fine-tuning task [6]. Recent work has shown that small models exhibit a fundamentally different learning capacity under fine-tuning compared to their larger counterparts [11], suggesting that interventions calibrated for large models may not transfer directly to SLMs.

This characterizes how certain layers can be more receptive to interventions towards certain tasks, especially when it comes to lighter and smaller language models.

III. METHODOLOGY

The methodology adopted in this paper follows a layer-wise intervention framework to systematically analyze how Representation Fine-Tuning (ReFT) affects performance across different transformer depths under resource-constrained settings. Each transformer layer is fine-tuned within the same set of hyperparameters like the low-rank dimension, batch size and other configuration variables. Because all the hyperparameters are the same except the layer index, this allows us to carry out direct comparison between the scores of the layers, their contribution to the score and how do they perform.

TABLE I: Hyperparameters for fine-tuning across all benchmark datasets.

Hyperparameter	SST-2	HellaSwag	MNLI	QQP
Batch Size	32	4	32	32
Learning Rate	6e-4	2.5e-4	2.5e-4	3e-4
Training Samples	1200	1200	5000	5000
Evaluation Samples	850	850	1000	1000
Optimizer	AdamW	AdamW	AdamW	AdamW
Max Length	256	512	256	128

The three models used in this study differ in both architecture and scale. Qwen3-0.6B [12] is a decoder-only language model with 28 transformer layers and approximately 0.6B parameters, notable for its use of Grouped Query Attention (GQA) which reduces memory overhead during inference. **Granite-350M** is also a decoder-only model, similarly structured with 28 layers but a smaller parameter count of roughly 350M. RoBERTa-Base [14] sits on the opposite end of the architectural spectrum: an encoder-only model with just 12 layers and approximately 125M parameters, distinguished by its bidirectional attention that allows every layer to attend to the full input sequence simultaneously. This last property makes RoBERTa fundamentally different from the two decoder models, and as will be seen in the results, leads to notably different layer-wise behavior under ReFT.

Evaluation is performed on a set of benchmarks on different nature of tasks like sentiment analysis, commonsense reasoning, and natural language inference, including SST-2 [15], HellaSwag [16], MNLI [18], and QQP [17]. For decoder-only models, HellaSwag and SST-2 are evaluated using a generative single-token classification setup, where the model is prompted with the input context and corresponding answer choices, and is required to generate one token from the label space A, B, C, D; accuracy is computed based on exact match with the ground-truth label. This formulation reframes multiple-choice evaluation as constrained token generation, enabling direct assessment of how ReFT influences decision-making behavior in constrained and regressive settings. In addition to accuracy and loss-based metrics, representational changes induced by intervention are analyzed using Singular Value Decomposition (SVD) to study similarity structure and rank

dynamics of learned intervention parameters across layers. All experiments are implemented using Google Colab and Kaggle environments, primarily utilizing NVIDIA T4 GPUs for training and evaluation.

IV. EXPERIMENTS

ReFT is evaluated across three models across four benchmark tasks. For each model, ReFT is applied independently to a single transformer layer at a time while keeping hyperparameters the same. The results are analyzed across layer-wise trends, training loss and gradient dynamics. This analysis potentially offers an explanation about whether ReFT is effective for lightweight models on tight dataset constraints or not.

A. Layer-wise Accuracy and Task Sensitivity

Table II summarizes the layer-wise accuracy statistics across all model-task combinations. A clear pattern emerges when comparing variance across tasks: sentiment classification (SST-2) exhibits substantially lower variance than commonsense reasoning (HellaSwag) for both decoder models, with Granite-350M on SST-2 recording a standard deviation of just 0.36% compared to 6.86% on HellaSwag. This suggests that for easier and well-defined tasks, ReFT interventions are largely not dependent on layer selection - the model’s representations are good enough at any layer for accurate classification. Conversely, the increased deviation on HellaSwag indicates that commonsense reasoning is considerably more sensitive to where in the network the intervention is applied, implying that task complexity is a primary moderator of layer selection sensitivity. Furthermore, one more observation can also be made: smaller models (Granite-350M and RoBERTa-Base) showed less deviation in their score than larger model (Qwen-0.6B) based on the layer, suggesting that larger models have richer representations, and have the better optimization surface for intervention compared to smaller models. Smaller architectures on the other hand maintain a more uniform representation across their depth.

The grouped accuracy statistics in Table III further reveal diverging trends across model families. For Qwen-0.6B on both tasks, late layers (20–27) show a consistent drop in mean accuracy relative to early and middle layers — most pronounced on HellaSwag, where late-layer mean accuracy falls to 31.07% compared to 47.85% in middle layers. Granite-350M, by contrast, exhibits a monotonic increase across layer groups on SST-2 (89.55% → 89.81% → 89.96%), suggesting that later representations in this model become progressively more amenable to intervention for simple classification. On HellaSwag, Granite peaks in the middle group (35.48%) before declining, a pattern consistent with the hypothesis that middle layers encode the richest representations for tougher tasks. The accuracy trend for RoBERTa-Base on both MNLI and QQP diverge from the hypothesis, where each layer is competent enough and leans into the same accuracy group. One possible explanation could be the architectural difference: RoBERTa base is masked, bidirectional language model where representations from past as well as the future are encoded and

saturated relatively equal for all the layers, making them less ‘specialized’ compared to decoder-only models, where future context is not involved.

There was one notable outlier: Layer 1 of Granite-350M on HellaSwag produced a validation accuracy of 0.0, excluded from the aggregate statistics in Table II. Further analysis of the loss reveals that intervention at this layer collapsed entirely to generating a single malformed token, resulting in an accuracy of 0%. Potential causes of this catastrophic intervention collapse could include poor initialized, unstable optimization at that depth. The layer is further inspected in the loss section to see its training behavior IV-B.

B. Training Dynamics and Gradient Behavior

Training dynamics across layers reveal patterns that both reinforce and extend the accuracy findings discussed above. Figures 2 and 3 present the per-layer training loss trajectories and gradient norm trends respectively across all model-task combinations. The training behavior of the layers suggests that the loss landscape is highly sensitive to the model architecture and the corresponding task.

For Granite-350M on SST-2, training loss curves across all layers converge smoothly where the final loss of all the layers is more or less the same. This similarity in nature is also visible in the low variance of the accuracies on SST-2 reported in Table II, further providing more evidence that simpler and well-defined tasks like SST-2 have more predictable and smoother loss landscape. Similarly, the gradient norms are also fairly consistent with no clear outliers. The overarching trends is that for simpler tasks, any layer is good enough to do the job for fine-tuning, reaching a fairly accurate solution at the end as the loss and gradient norm trends are overall the same on Granite-350M.

The behavior on HellaSwag is quite variable. Gradient norms for both Granite-350M and Qwen-0.6B on HellaSwag are more noisy, with sharp spikes persisting well beyond the initial steps rather than stabilizing. This behavior is consistent with higher variance observed in HellaSwag across layers, providing a further account of task sensitivity: ReFT interventions on reasoning tasks is noisier, making the effect of layer choice more vital.

When it comes to gradient norm plots of Qwen-0.6B, one pattern that can be observed is: certain layers cluster together in terms of gradient norms and loss trajectory, while other layers do not showcase the same behavior. This clustering behavior is less prominent in Granite-350M, where clusters are harder to identify. One plausible interpretation of this behavior is that despite being small, Qwen-0.6B exhibits more specialized behavior across its layers. This finding can be projected to larger models, where this behavior is possibly more prominent. Granite-350M, being smaller still, may lack the capacity for such specialization, with all layers accounting to more similar representations, doing the heavy lifting.

Late layers across both decoder models tend to exhibit higher training loss at convergence relative to early and middle layers, paired with relatively lower gradient norm magnitudes.

TABLE II: Layer-wise ReFT Accuracy% Statistics per Model and Task

Model	Task	Mean Accuracy	Standard Deviation	Best Accuracy	$\text{Acc}_{\max} - \text{Acc}_{\min}$
Qwen-0.6B	SST-2	88.54	1.68	90.47	6
Qwen-0.6B	HellaSwag	42.98	8.47	49.29	22
Granite-350M	SST-2	89.76	0.366	90.35	1.53
Granite-350M	HellaSwag	33.78	6.86	38.94	7.65 [†]
RoBERTa-Base	MNLI	44.01	2.94	50.20	9.7
RoBERTa-Base	QQP	73.91	1.85	77.00	6

[†]Layer 1 (accuracy = 0.0) excluded due to generation collapse; see Section IV-A.

TABLE III: Mean Accuracy% by Layer Group per Model and Task

Model	Task	Early Layers	Middle Layers	Late Layers
Qwen-0.6B	SST-2	89.34 (0–9)	89.40 (10–19)	86.47 (20–27)
Qwen-0.6B	HellaSwag	47.64 (0–9)	47.85 (10–19)	31.07 (20–27)
Granite-350M	SST-2	89.55 (0–9)	89.81 (10–19)	89.96 (20–27)
Granite-350M	HellaSwag	32.31 (0–9)	35.48 (10–19)	33.50 (20–27)
RoBERTa-Base	MNLI	41.50 (0–2)	43.33 (3–5)	44.58 (6–9)
RoBERTa-Base	QQP	72.60 (0–2)	72.70 (3–5)	74.75 (6–9)




Fig. 1: Layer-wise ReFT validation accuracy across all model-task combinations. Dashed vertical lines delineate early, middle, and late layer regions. Layer 1 of Granite-350M on HellaSwag (accuracy = 0.0) is visible in (d) and discussed in Section IV-A.

This combination of high loss and weaker gradient norms suggests that late layers are less receptive to ReFT interventions. One plausible interpretation is that later layers are geared more towards structuring the output instead of semantic richness, making them weaker candidates for ReFT interventions. This is particularly evident in Qwen-0.6B on SST-2, where one layer remains at an anomalously high loss throughout the entire training run without converging, an outlier in Figure 2.

Inspecting the gradient norm and loss trajectory of Layer 1 of Granite-350M, the layer that catastrophically collapsed, we can observe that the loss spiked during training and it

never recovered for that specific layer. Similarly, the gradient norms for this layer are correspondingly erratic. Prediction analysis confirms the consequence: the model outputs a single malformed token for every example, producing 0.0 validation accuracy. This underscores the practical importance of monitoring training loss trajectories when applying ReFT to small models, as catastrophic collapse can occur silently from an accuracy standpoint if evaluation is not performed at each layer independently.

RoBERTa-Base presents a qualitatively different optimization profile. Training loss curves are noisier step-to-step than




Fig. 2: Per-layer ReFT training loss trajectories across all model-task combinations. Each line corresponds to a single transformer layer. The anomalous spike in (b) corresponds to Layer 1 of Granite-350M on HellaSwag, which diverges and never recovers, consistent with the generation collapse discussed in Section IV-A. Darker shades correspond to earlier layers and lighter shades correspond to later layers.

those of the decoder models on SST-2, reflecting the larger training set and longer training duration, but show a clear downward trend across all layers without catastrophic divergence. Gradient norms remain active and relatively uniform across layers throughout training, with no strong evidence of the late-layer reluctance pattern seen in the decoder models. This is consistent with the accuracy finding that RoBERTa’s later layers are in fact its most productive intervention points — if later layers were resistant to optimization, one would expect elevated loss or diminished gradients there, but neither is observed. The encoder architecture’s bidirectional attention may contribute to a more uniform representational structure across depths, making all layers comparably trainable while still allowing later layers to encode richer semantic content relevant to inference tasks like MNLI and QQP.

C. Intervention Weight Analysis


Beyond accuracy and training dynamics, the learned intervention weights themselves offer a window into what ReFT captures at each layer. For each trained intervention, three parameter sets are available for analysis: the learned weight matrix W , the bias vector b , and the rotation matrix R , as we are applying the LoReFT. We analyze these through several lenses — Frobenius norm as a measure of total intervention magnitude, effective rank as a measure of spectral diversity,

and deviation from orthogonality for the rotation matrix — and correlate these against per-layer validation accuracy using Pearson correlation coefficients.


Across all layers and both decoder models, the rotation matrix R behaves identically. Its deviation from orthogonality stays near zero, proving that R consistently maintains its geometric structure as a near-perfect orthonormal basis at every intervention point. Furthermore, the effective rank of R remains stable across layers and shows no meaningful correlation with model accuracy under any configuration.

The weight matrix W is more significant. The Frobenius norm of W correlates strongly with accuracy for Qwen-0.6B on HellaSwag, with a Pearson coefficient of $r = 0.792$, the highest observed across all model-task combinations. This indicates that for reasoning tasks in this model, layers where the intervention learns larger weight magnitudes tend to produce better accuracy outcomes — the intervention is doing more work, and that work is productive.


The same relationship does not hold uniformly across configurations. For SST-2, the correlation between weight norm and accuracy is negative for both Qwen-0.6B $r = -0.491$ and Granite-350M $r = -0.352$, suggesting that on simpler tasks, larger intervention magnitude may reflect overfitting to the training signal rather than meaningful representational




(a) Granite-350M on SST-2



(b) Granite-350M on HellaSwag



(c) Qwen-0.6B on SST-2



(d) Qwen-0.6B on HellaSwag

Fig. 3: Gradient norm trends across layers during ReFT training for decoder models. Log scale is used on the y-axis to accommodate the wide dynamic range. Higher volatility in HellaSwag configurations relative to SST-2 is consistent with the elevated accuracy variance reported in Table II. Darker shades correspond to earlier layers and lighter shades correspond to later layers.

steering. This asymmetry across tasks is itself a meaningful finding: the relationship between intervention magnitude and performance is not universal but steered by task complexity.

The strong positive correlation between weight norm and accuracy on Qwen-HellaSwag is further reinforced when considered alongside the accuracy collapse at layers 21–27 in that configuration. These late layers, which produce the lowest accuracy scores, also correspond to interventions with notably lower weight norms.

D. Token Prediction

As we are using single token generation for classification, we can analyze the token outputs generated by these models. For Qwen-0.6B, from layer 0 to 20, the token generation is more or less uniform, where all four options are generated with no biased generation of a single option, and these were the layers before the we lost the accuracy gains. For example, layer 20 generated predictions as 28.82% (A), 25.76% (B), 23.29% (C), and 22.12% (D), closely matching the distribution

of labels of the validation dataset. But after layer 20, the trend changes, where the token generation becomes heavily biased towards a specific option. For instance, one of the worst performing layer, layer 21, right after layer 20, generated 79.41% of all predictions as label A. This sharp transition potentially explains how there are certain 'structural boundaries' within the layers depending on the task, where ReFT greatly loses its effectiveness. And this token generation bias only worsens from now on, such as label A is generated as highly as 92.24% at layer 24.

A uniquely different and complete collapse is observed at Layer 1 of Granite-350M on HellaSwag. Here the model does not fixate on a valid answer label but instead generates a single malformed token for 100% of evaluated examples, producing 0.0 validation accuracy. Unlike the Qwen late-layer collapse, the Granite Layer 1 failure is total and immediate, pointing to a fundamentally different failure mode: bad initialization-level divergence at a layer too early in the network, which is further evident observing the loss and gradient norms of this layer.




Fig. 4: Model intervention capacity and structural metrics correlation with accuracy




Fig. 5: Layer 20 vs Layer 21 Token Prediction Statistics, highlighting collapse in token prediction output

V. DISCUSSION

The results across all three models and four tasks promote an important finding: layer selection for ReFT intervention is quite sensitive for task. For sentiment classification, ReFT is robust and layer-agnostic for the most part. You can get quite similar results on any layer when dealing with simple and clearly defined tasks, such as SST-2, and the layers can reach a workable solution within comparable number of epochs. However, when it comes to more complex tasks like commonsense reasoning, the selection becomes pivotal. Consistently, the middle layer interventions have performed better because of rich semantic representations they hold, while the later layers are more geared for output token generation and positional encoding. This is evident from the loss landscape of late layers, showcasing higher loss with lower gradient norms. We also observed the disparity of layers between relatively bigger

and smaller models, where bigger models (like Qwen-0.6B) with ReFT showed better score but with more variance, while smaller models (like Granite-350M and RoBERTa) had lower variance in the accuracies but also performed less, with lower overall scores.

Encoder and decoder models also differ in how they respond to layer-level interventions. RoBERTa-Base consistently benefits from interventions at later layers on both MNLI and QQP — the opposite of what is observed in decoder models. This reversal is unlikely to be explained by model size alone, since Granite-350M is comparably sized yet follows the decoder-typical pattern of degradation at later layers. A more plausible explanation lies in the architectural difference between bidirectional and causal attention. Because encoder models attend to the full input context at every layer, representations are progressively refined across depth, and later layers may encode more abstract semantic information that is more compatible with representation-level steering. In causal decoder models, later layers are more strongly shaped by the autoregressive objective, encoding signals that are less amenable to geometric intervention. Notably, the original LoReFT paper intervenes across all layers simultaneously when evaluating on GLUE and does not analyze individual layer depth [6], so the layer-wise pattern observed here is not directly reported in prior work.

The two collapse cases observed — more tokens being a single option in Qwen-0.6B at Layer 21 on HellaSwag and malformed token generation in Granite-350M at Layer 1 — represent distinct failure modes with different practical implications. The Qwen collapse is more of a specialized threshold effect, where the model starts focusing on the generation of the tokens instead of calibrating the output after a certain layer, beyond which the intervention can no longer support the task. In practice, this means that for HellaSwag on Qwen-0.6B, layers beyond 20 can be safely excluded from the search space. The Granite Layer 1 collapse is different in nature, reflecting an optimization instability at a shallow depth that results in total divergence. This is harder to anticipate, but it is detectable early — the divergence appears within the first few training iterations, so full training is not required to identify it.

The intervention weight analysis provides additional insight into where and why ReFT succeeds. The rotation matrix R remains near-perfectly orthogonal across layers and contributes little to accuracy variance, suggesting that the geometric subspace constraint in LoReFT is less critical than the learned transformation within it. This implies that simplifying or fixing R in future ReFT variants may have little practical cost, while allocating more capacity to the weight matrix W could be more beneficial. The finding that $\|W\|_F$ correlates positively with accuracy on HellaSwag but negatively on SST-2 further suggests that intervention magnitude is not a reliable proxy for intervention quality — its effect depends on the task.

VI. LIMITATIONS

Because of limited hardware (single NVIDIA T4 GPU), the training sets samples used are small. While this reflects the

low-resource setting which the paper focuses on, it limits the generalizability of findings to standard fine-tuning regimes, where larger datasets are utilized with more compute. It is plausible that bigger dataset may make the late layers more effective and generalized which may yield better score from them. Secondly, this paper focuses on intervening one layer at a time. While it has demonstrated which layers and layer groups achieve better results, multi-layer interventions have not been demonstrated, which may yield potentially even better results, but this would also require more compute and the search space will be substantial.

Furthermore, three models were tested in this paper: Qwen-0.6B, Granite-350M, RoBERTa-Base. While some trends were observed depending on the model architecture and parameter space, evaluating larger models like Phi-3-mini or Gemma-2B would make the results more generalized and concrete to see if the similar trends emerge in larger models, with more specialized layers, better accuracies and whether the middle layers still emerge as the better candidate for ReFT interventions.

REFERENCES

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, "Attention is all you need," in *Proc. Advances in Neural Information Processing Systems (NeurIPS)*, 2017, pp. 6000–6010.
- [2] T. B. Brown, B. Mann, N. Ryder, M. Subbiah, J. D. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell, S. Agarwal, A. Herbert-Voss, G. Krueger, T. Henighan, R. Child, A. Ramesh, D. Ziegler, J. Wu, C. Winter, C. Hesse, M. Chen, E. Sigler, M. Litwin, S. Gray, B. Chess, J. Clark, C. Berner, S. McCandlish, A. Radford, I. Sutskever, and D. Amodei, "Language models are few-shot learners," in *Proc. Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [3] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen, "LoRA: Low-rank adaptation of large language models," in *Proc. Int. Conf. on Learning Representations (ICLR)*, 2022.
- [4] N. Hounsby, A. Giurgiu, S. Jastrzebski, B. Morrone, Q. de Laroussilhe, A. Gesmundo, M. Attariyan, and S. Gelly, "Parameter-efficient transfer learning for NLP," in *Proc. Int. Conf. on Machine Learning (ICML)*, 2019, pp. 2790–2799.
- [5] X. L. Li and P. Liang, "Prefix-tuning: Optimizing continuous prompts for generation," in *Proc. Annual Meeting of the Association for Computational Linguistics (ACL)*, 2021.
- [6] Z. Wu, A. Arora, Z. Wang, A. Geiger, D. Jurafsky, C. D. Manning, and C. Potts, "ReFT: Representation Finetuning for Language Models," *arXiv preprint arXiv:2404.03592*, 2024.
- [7] A. Geiger, Z. Wu, C. Potts, T. Icard, and N. D. Goodman, "Finding Alignments Between Interpretable Causal Variables and Distributed Neural Representations," *arXiv preprint arXiv:2303.02536*, 2023.
- [8] W. Gurnee and M. Tegmark, "Language Models Represent Space and Time," *arXiv preprint arXiv:2310.02207*, 2023.
- [9] J. Engels, E. J. Michaud, I. Liao, W. Gurnee, and M. Tegmark, "Not All Language Model Features Are Linear," *arXiv preprint arXiv:2405.14860*, 2024.
- [10] X. Song, K. Wang, P. Li, L. Yin, and S. Liu, "Demystifying the Roles of LLM Layers in Retrieval, Knowledge, and Reasoning," *arXiv preprint arXiv:2510.02091*, 2025.
- [11] Y. Li, X. Yue, Z. Xu, F. Jiang, L. Niu, B. Y. Lin, B. Ramasubramanian, and R. Poovendran, "Small Models Struggle to Learn from Strong Reasoners," *arXiv preprint arXiv:2502.12143*, 2025.
- [12] A. Yang, A. Li, B. Yang, et al., "Qwen3 technical report," *arXiv preprint arXiv:2505.09388*, 2025.
- [13] Gemma Team, "Gemma 3 technical report," *arXiv preprint arXiv:2503.19786*, 2025.
- [14] Y. Liu, M. Ott, N. Goyal, et al., "RoBERTa: A robustly optimized BERT pretraining approach," *arXiv preprint arXiv:1907.11692*, 2019.
- [15] R. Socher, A. Perelygin, J. Wu, J. Chuang, C. D. Manning, A. Ng, and C. Potts, "Recursive deep models for semantic compositionality over a sentiment treebank," in *Proc. Conf. on Empirical Methods in Natural Language Processing (EMNLP)*, 2013.
- [16] R. Zellers, A. Holtzman, Y. Bisk, A. Farhadi, and Y. Choi, "Hel-laSwag: Can a machine really finish your sentence?" *arXiv preprint arXiv:1905.07830*, 2019.
- [17] S. Iyer, N. Dandekar, and K. Csernai, "First Quora dataset release: Question pairs," Quora, 2017.
- [18] A. Williams, N. Nangia, and S. Bowman, "A broad-coverage challenge corpus for sentence understanding through inference," *arXiv preprint arXiv:1704.05426*, 2017.